



# Deep Generative Modeling - part 2 -

Hana Sebia

[hana.sebia@inria.fr](mailto:hana.sebia@inria.fr)

Decemeber, 2023

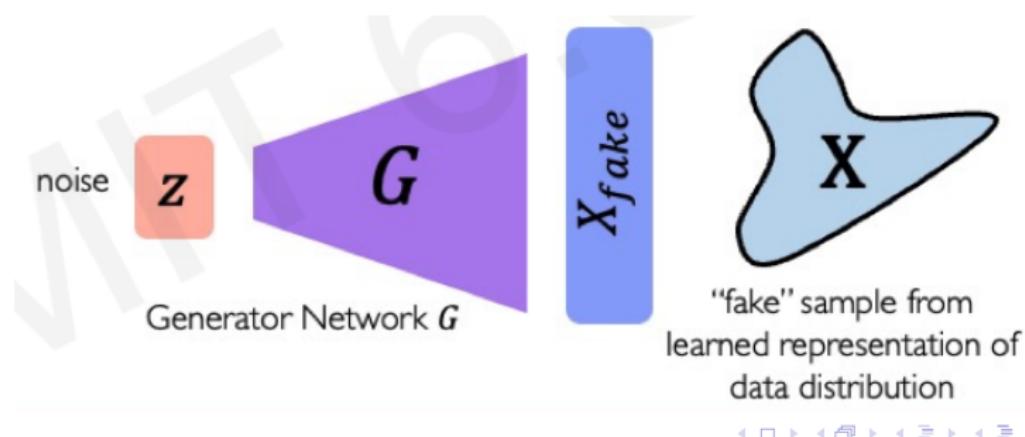
AIstroSight, Inria, Lyon Center

What if we just want to sample ?

**Idea** : don't explicitly model density, and just sample to generate new instances

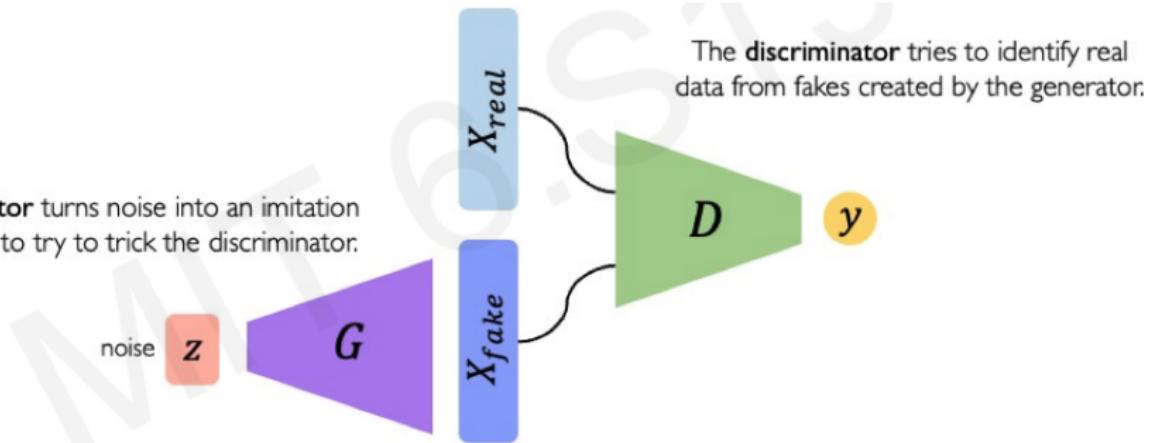
**Problem** : sample from complex distribution - can't do this directly !

**Solution** : sample from something simple, learn a transformation to the data distribution.



**Generative Adversarial Networks (GANs)** are a way to make generative model by having two neural networks compete with each other.

The **generator** turns noise into an imitation of the data to try to trick the discriminator.



**Generator** starts from noise to try to create an imitation of the data

Generator



**Discriminator** looks at both real data and fake data created by the generator

Discriminator



Generator



**Discriminator** tries to predict what's real and what's fake

Discriminator



Generator

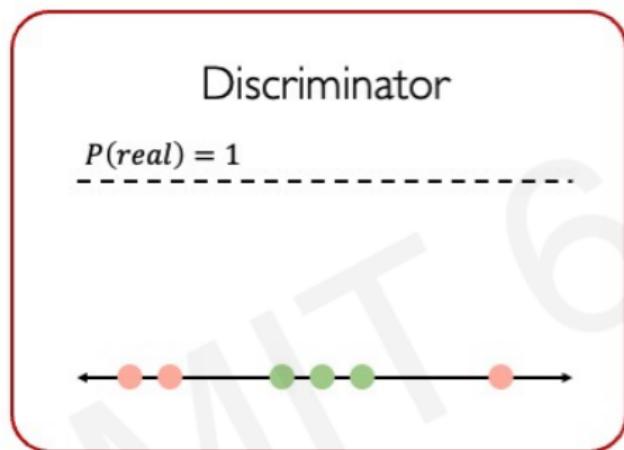


Real data



Fake data

Discriminator tries to predict what's real and what's fake



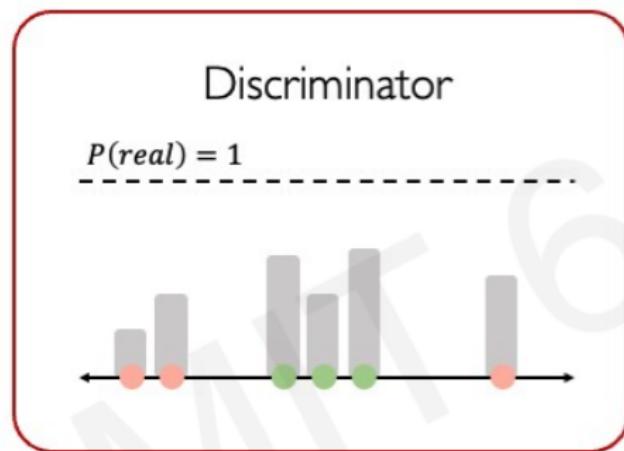
Generator



Real data

Fake data

Discriminator tries to predict what's real and what's fake



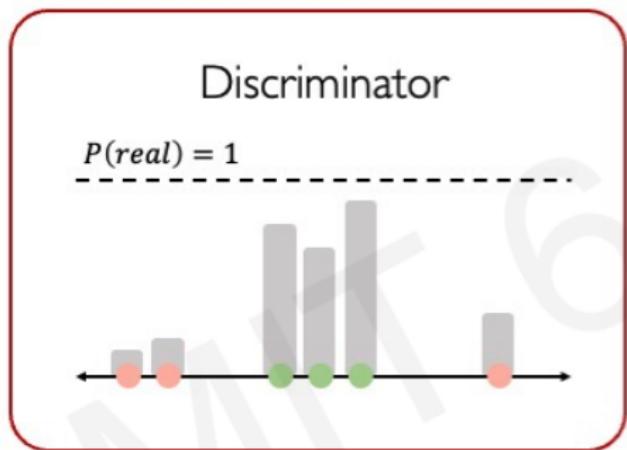
Generator



Real data

Fake data

Discriminator tries to predict what's real and what's fake



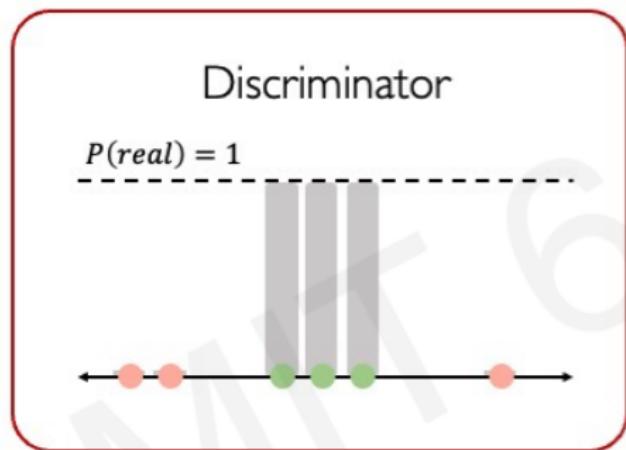
Generator



Real data

Fake data

Discriminator tries to predict what's real and what's fake



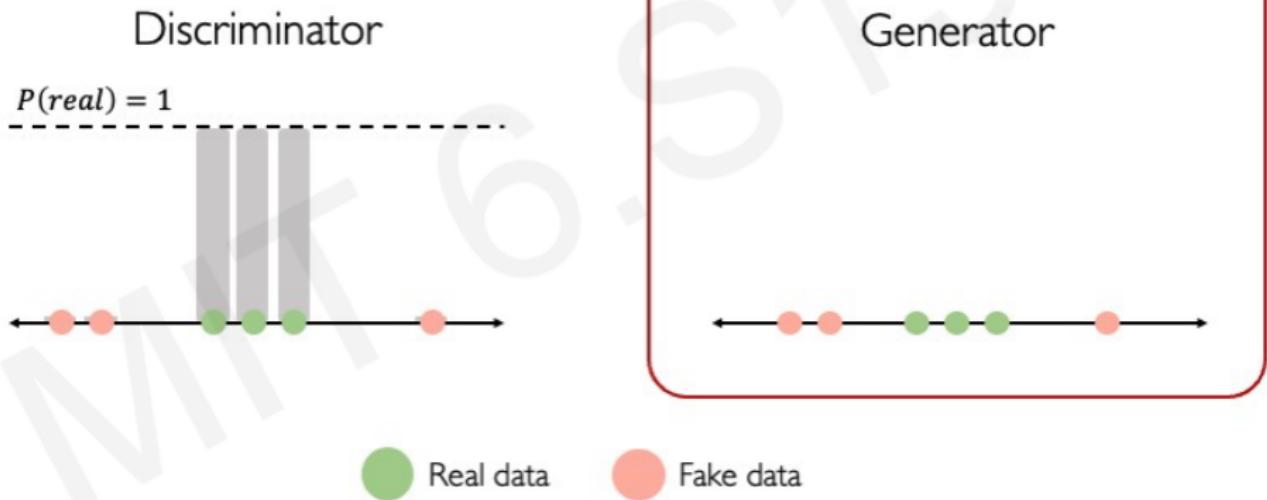
Real data

Fake data

**Generator** tries to improve its imitation of the data



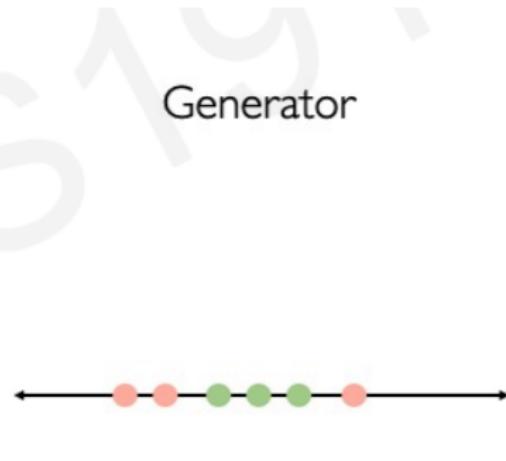
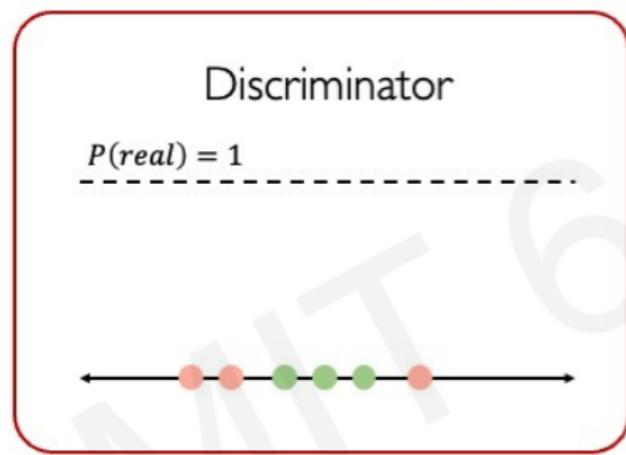
**Generator** tries to improve its imitation of the data



**Generator** tries to improve its imitation of the data



Discriminator tries to predict what's real and what's fake

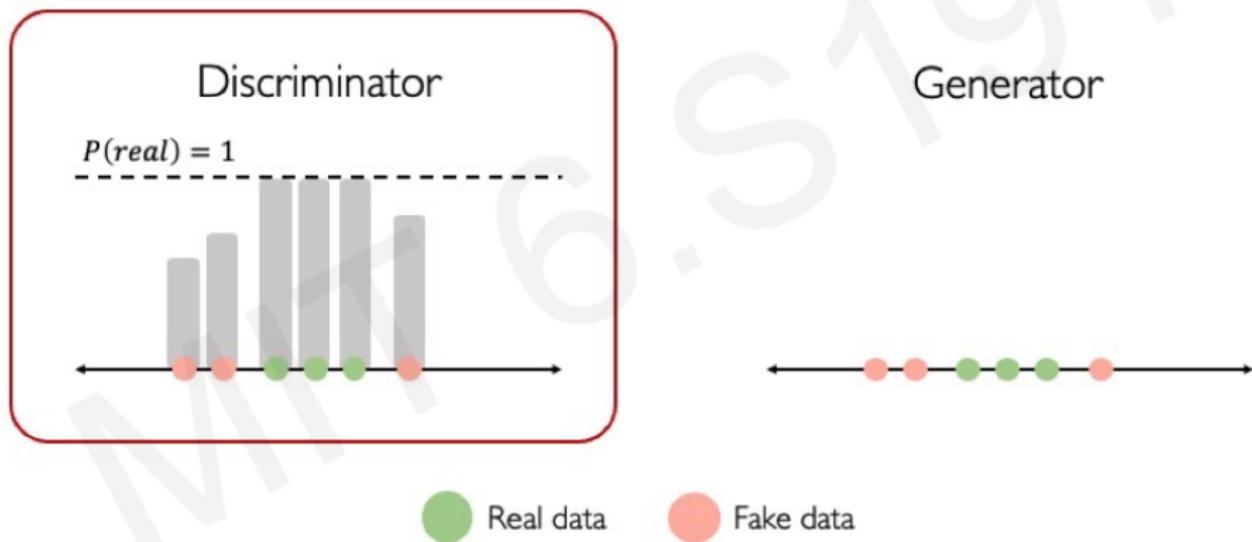


Real data

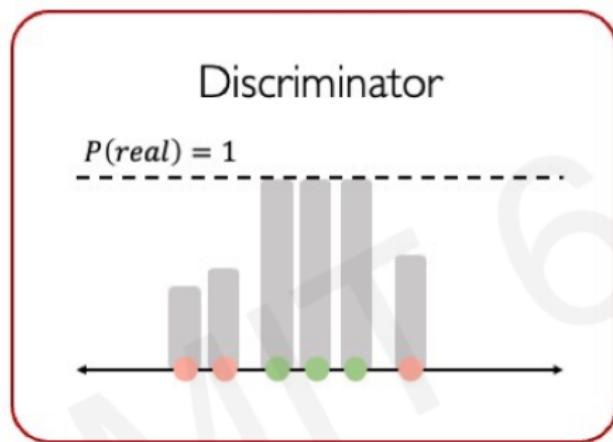


Fake data

**Discriminator** tries to predict what's real and what's fake



**Discriminator** tries to predict what's real and what's fake

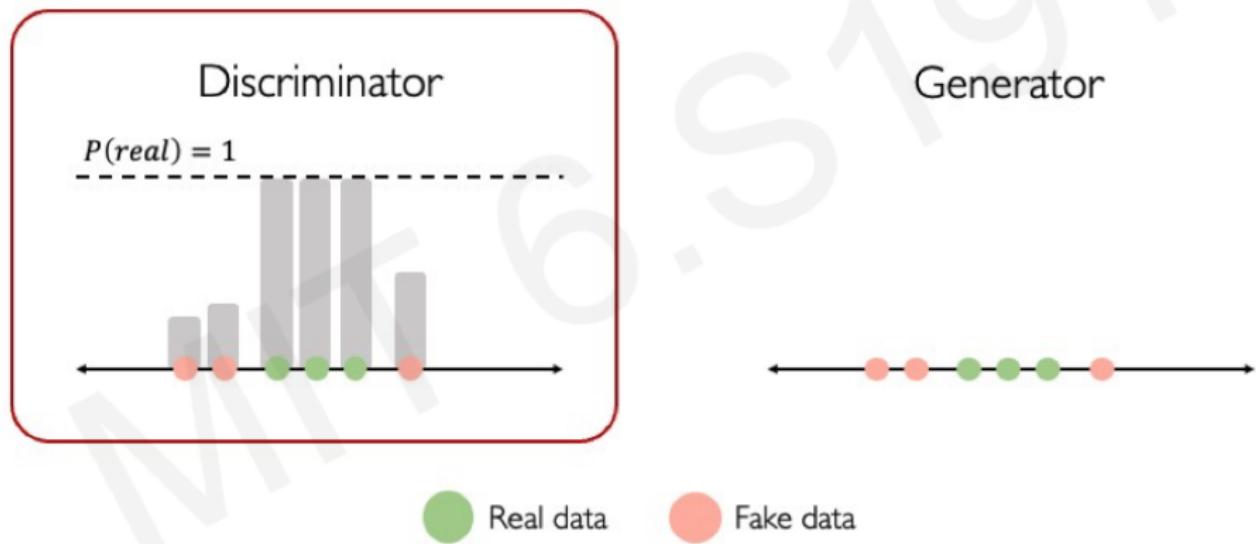


Generator



Real data   Fake data

**Discriminator** tries to predict what's real and what's fake



**Generator** tries to improve its imitation of the data



**Generator** tries to improve its imitation of the data



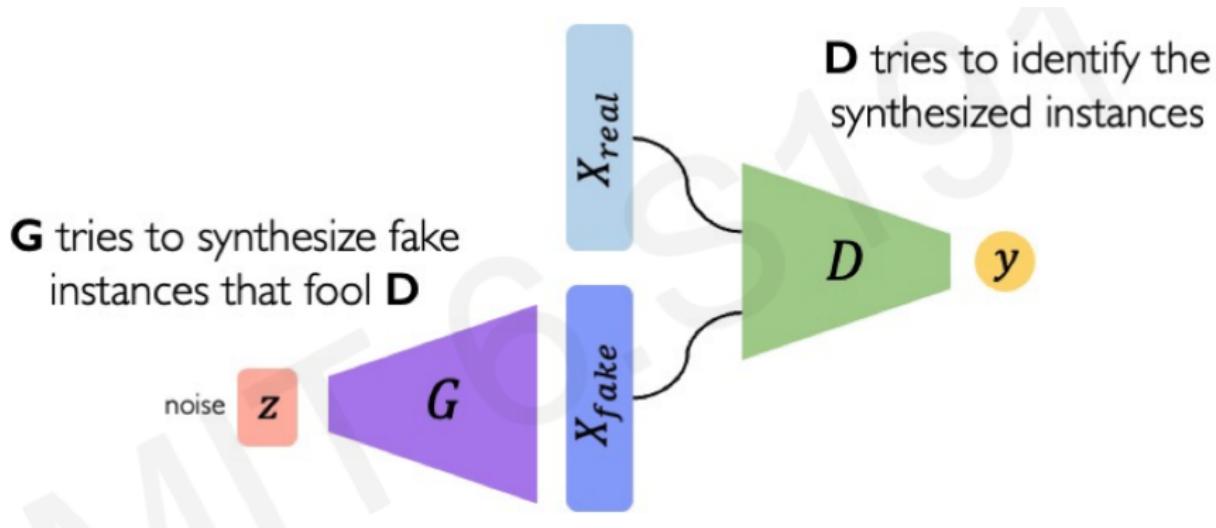
**Generator** tries to improve its imitation of the data



**Discriminator** tries to identify real data from fakes created by the generator

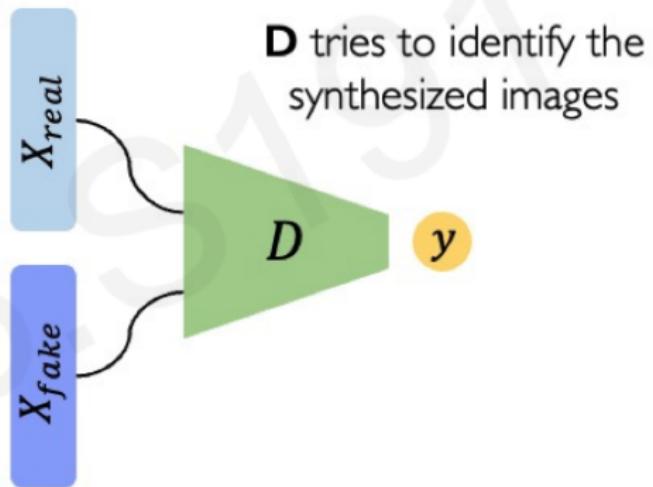
**Generator** tries to create imitations of the data to trick the discriminator





**Training:** adversarial objectives for **D** and **G**  
**Global Optimum:** **G** reproduces the true data distribution

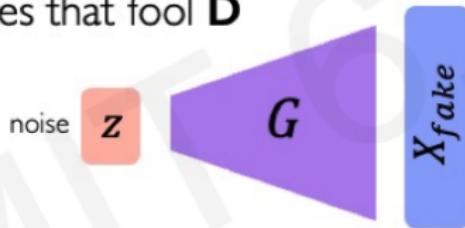
# Training GANs : loss function



$$\arg \max_D \mathbb{E}_{\mathbf{z}, \mathbf{x}} [ \frac{\log D(G(\mathbf{z}))}{\text{Fake}} + \frac{\log (1 - D(\mathbf{x}))}{\text{Real}} ]$$

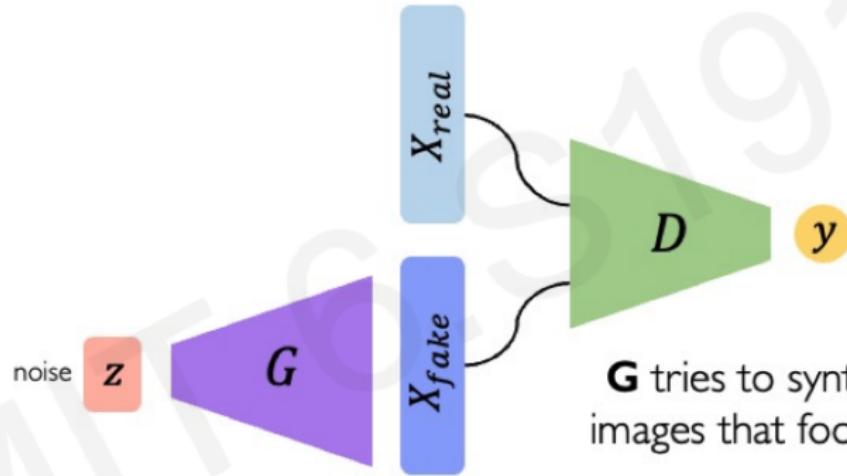
# Training GANs : loss function

**G** tries to synthesize fake images that fool **D**



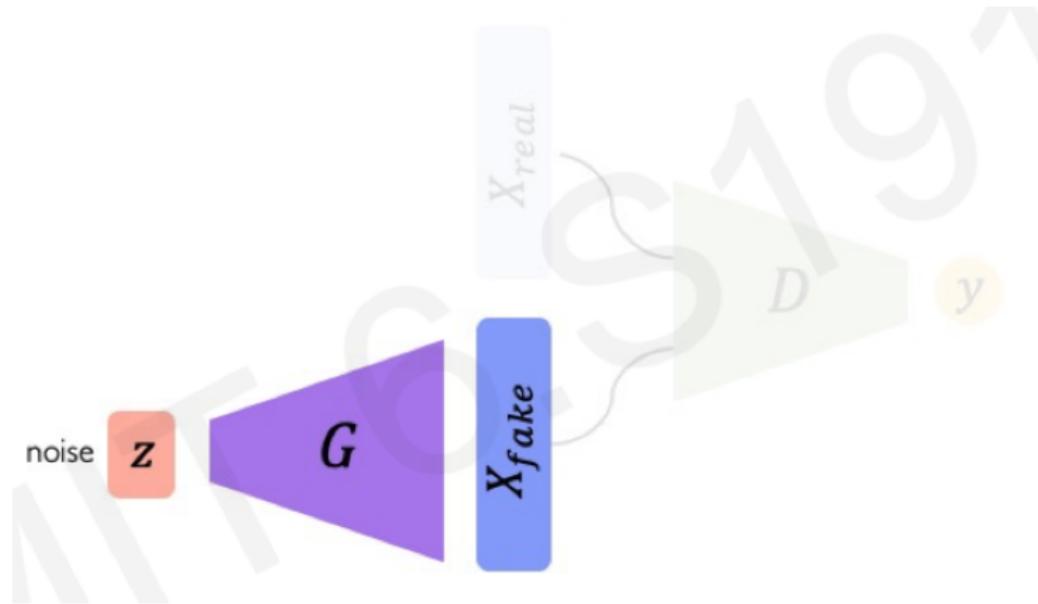
$$\arg \min_G \mathbb{E}_{\mathbf{z}, \mathbf{x}} [ \log D(G(\mathbf{z})) + \log (1 - D(\mathbf{x})) ]$$

# Training GANs : loss function



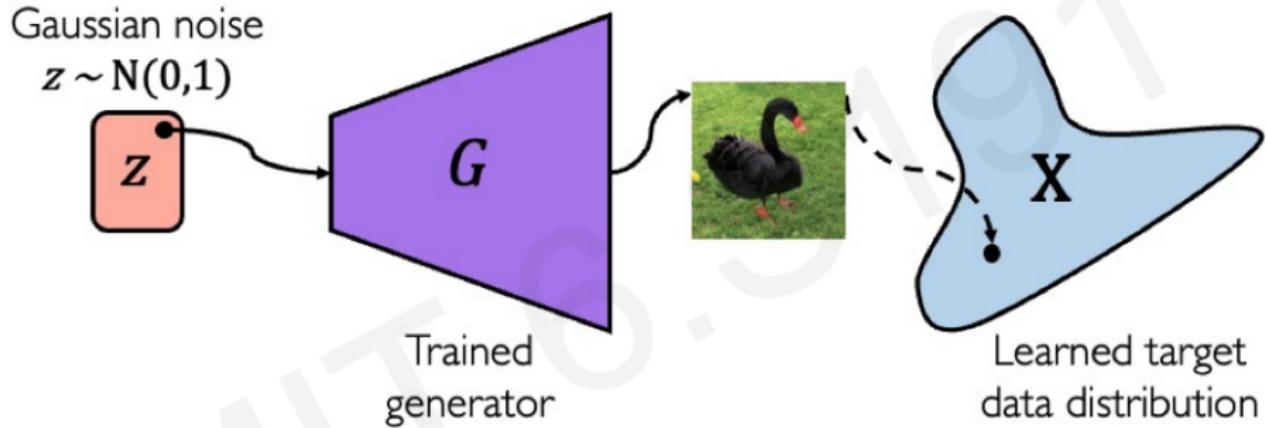
$$\arg \min_G \max_D \mathbb{E}_{\mathbf{z}, \mathbf{x}} [ \log D(G(\mathbf{z})) + \log (1 - D(\mathbf{x})) ]$$

# Generating new data with GANs



After training , use the generator network to create **new data** that's never been seen before

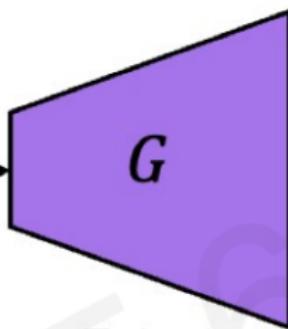
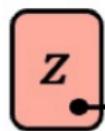
# GANS are distribution transformers



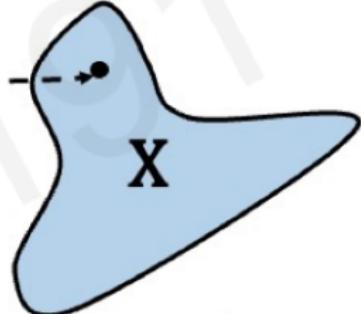
# GANS are distribution transformers

Gaussian noise

$$z \sim N(0,1)$$



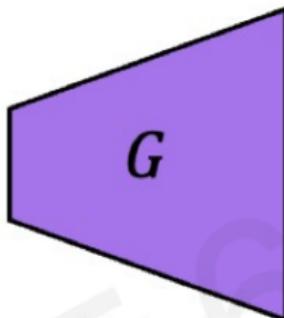
Trained  
generator



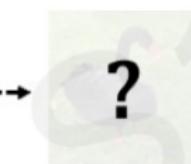
Learned target  
data distribution

# GANS are distribution transformers

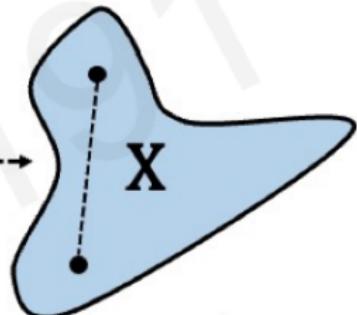
Gaussian noise  
 $z \sim N(0,1)$



Trained  
generator



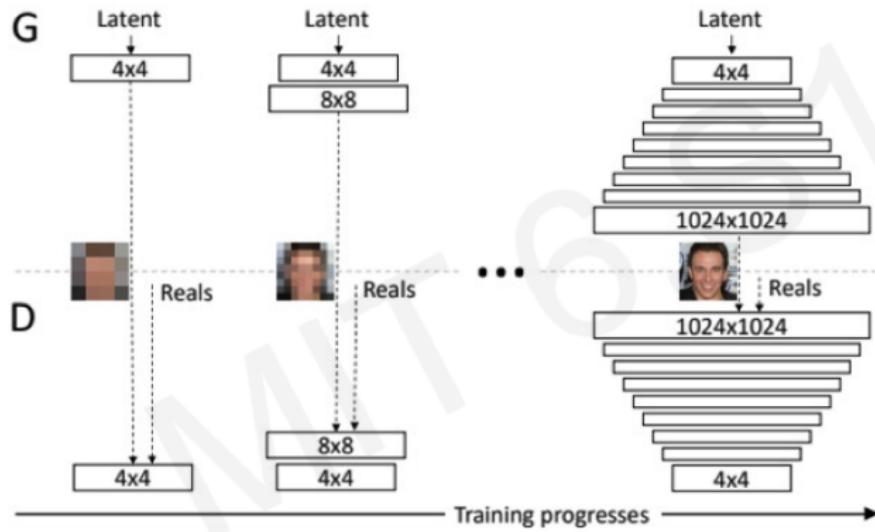
X



Learned target  
data distribution



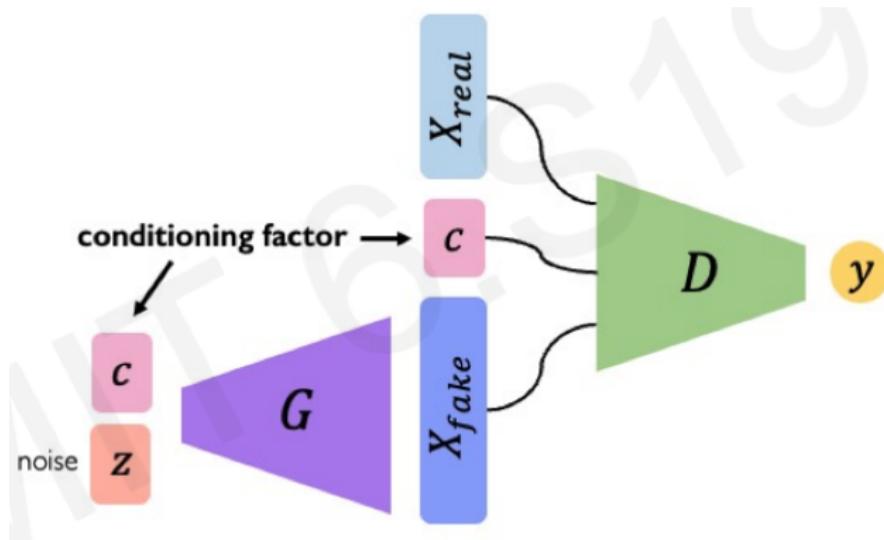
## Progressive growing of GANs



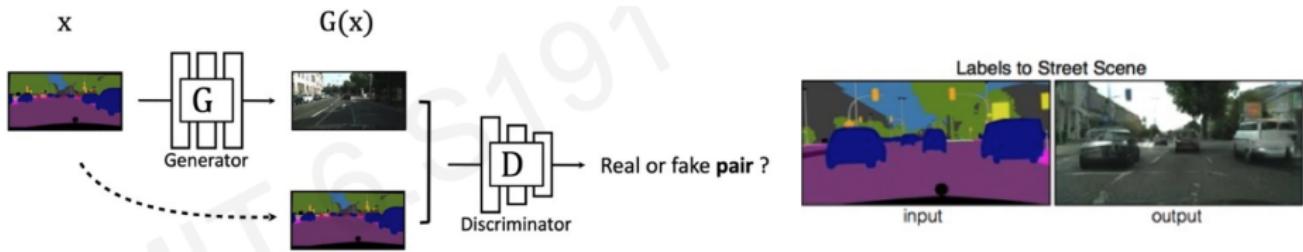
## Progressive growing of GANs : results



What if we want to control the nature of the output, by conditioning on a label ?



# Pix2pix: Paired translation

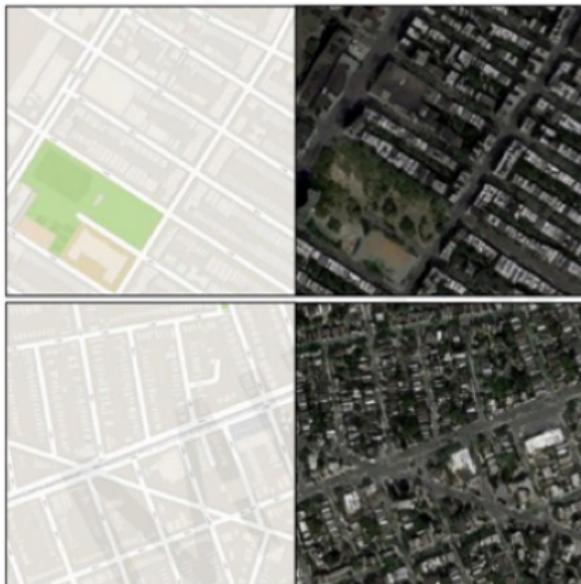


The discriminator, D, classifies between fake and real **pairs**.

The generator, G, learns to fool the discriminator.

# Paired Translation : results

Map → Aerial View



input

output

Aerial View → Map

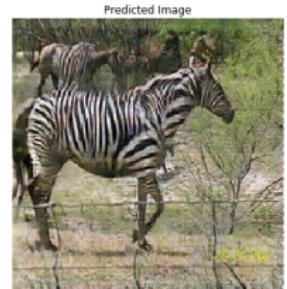
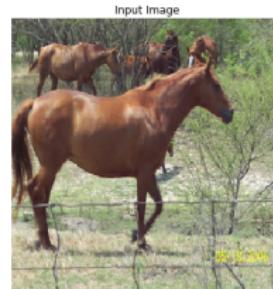
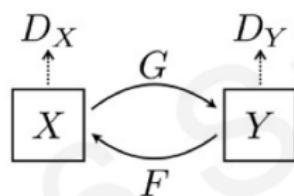


input

output

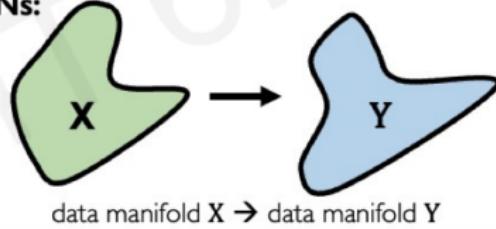
# CycleGAN: domain transformation

CycleGAN learns transformations across domains with unpaired data



# Distribution Transformation

CycleGANs:



data manifold X → data manifold Y

GANs:

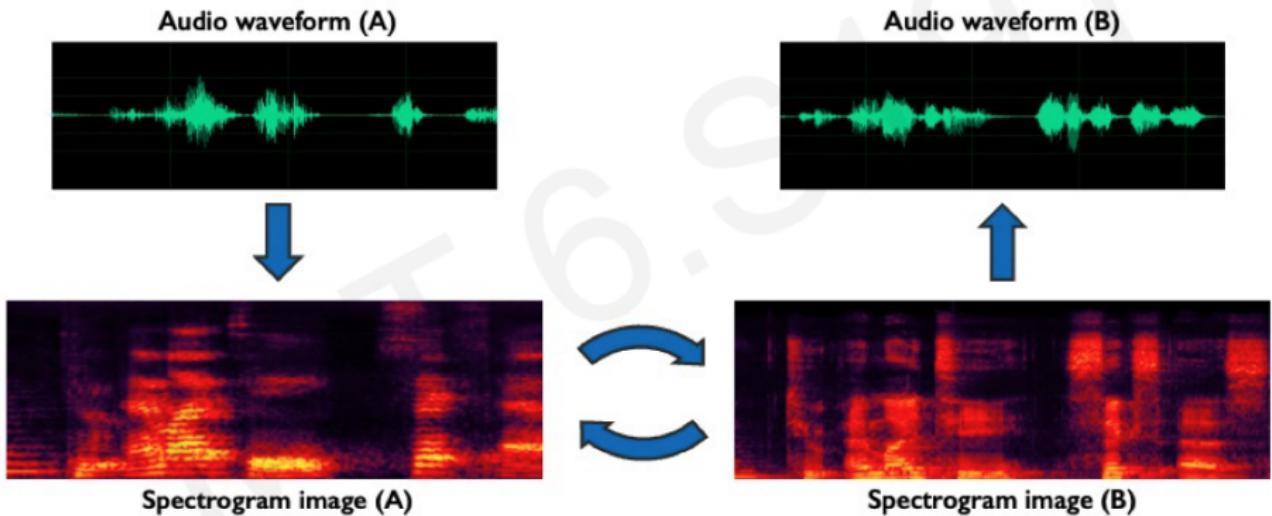
Gaussian noise  
 $z \sim N(0,1)$



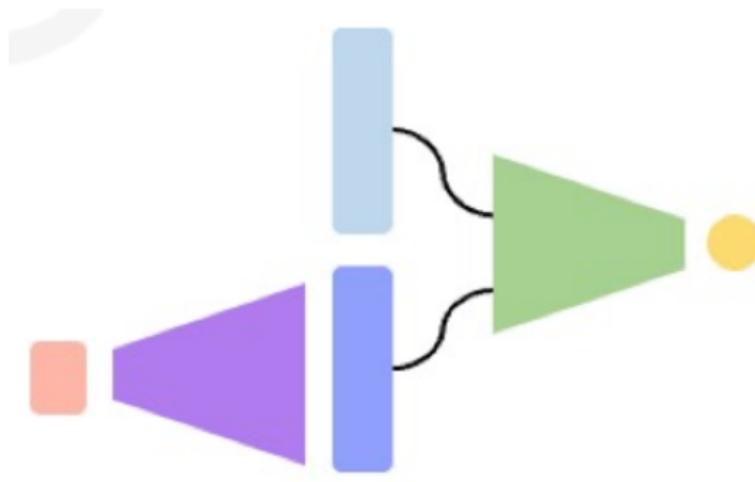
Gaussian noise → target data manifold

# CycleGAN : Transforming Speech

Inria



- ① Competing generator and discriminator networks
- ② Both networks can be trained using only backpropagation
- ③ Unsupervised training (no need for labels!)
- ④ Generating data that looks similar to original data
- ⑤ Hard to train



# Trilemma for GANs

- **High-fidelity samples** : sample quality is better (less blurry) than for VAE !
- **Low diversity samples** : also called **mode collapse**, this means the generator only learns to create a subset specialized in fooling discriminator.
- **Fast Sampling** : require just one pass of the generator network to turn noise into a realistic data sample

